

Assignment 3 : Gradient Descent

Aryaman Gautam

J001

```
In [146... %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Univariate Linear Regression

```
In [147... data = pd.read_csv("ex1data1.txt", header = None)
```

```
In [148... data.head()
```

```
Out[148...
      0      1
0  6.1101  17.5920
1  5.5277   9.1302
2  8.5186  13.6620
3  7.0032  11.8540
4  5.8598   6.8233
```

```
In [149... data.describe()
```

Out[149...

	0	1
count	97.000000	97.000000
mean	8.159800	5.839135
std	3.869884	5.510262
min	5.026900	-2.680700
25%	5.707700	1.986900
50%	6.589400	4.562300
75%	8.578100	7.046700
max	22.203000	24.147000

In [150...

```
data.columns = ['Profit', 'Population']
```

In [151...

```
plt.scatter(data['Population'], data['Profit'])  
plt.xticks(np.arange(5,30,step=5))  
plt.yticks(np.arange(-5,30,step=5))  
plt.xlabel('Population (in 10,000s)')  
plt.ylabel('Profit (in 10,000$)')  
plt.title('Profit vs Population')
```

Out[151...

Text(0.5, 1.0, 'Profit vs Population')



Cost Function $J(\theta)$

```
In [152... def computeCost(X,y,theta):
    """
    Take in a numpy array X,y,theta and get cost function using theta as parameter in a linear regression model
    """
    m = len(y)
    predictions = X.dot(theta)
    square_err = (predictions - y)**2

    return 1/(m)*np.sum(square_err)
```

```
In [153... data['x0'] = 1 #new column added
```

```
In [154... data.values
```

```
Out[154... array([[ 6.1101 , 17.592  ,  1.    ],
       [ 5.5277 ,  9.1302 ,  1.    ],
       [ 8.5186 , 13.662  ,  1.    ],
```

```
[ 7.0032 , 11.854 , 1. ],
[ 5.8598 , 6.8233 , 1. ],
[ 8.3829 , 11.886 , 1. ],
[ 7.4764 , 4.3483 , 1. ],
[ 8.5781 , 12. , 1. ],
[ 6.4862 , 6.5987 , 1. ],
[ 5.0546 , 3.8166 , 1. ],
[ 5.7107 , 3.2522 , 1. ],
[14.164 , 15.505 , 1. ],
[ 5.734 , 3.1551 , 1. ],
[ 8.4084 , 7.2258 , 1. ],
[ 5.6407 , 0.71618, 1. ],
[ 5.3794 , 3.5129 , 1. ],
[ 6.3654 , 5.3048 , 1. ],
[ 5.1301 , 0.56077, 1. ],
[ 6.4296 , 3.6518 , 1. ],
[ 7.0708 , 5.3893 , 1. ],
[ 6.1891 , 3.1386 , 1. ],
[20.27 , 21.767 , 1. ],
[ 5.4901 , 4.263 , 1. ],
[ 6.3261 , 5.1875 , 1. ],
[ 5.5649 , 3.0825 , 1. ],
[18.945 , 22.638 , 1. ],
[12.828 , 13.501 , 1. ],
[10.957 , 7.0467 , 1. ],
[13.176 , 14.692 , 1. ],
[22.203 , 24.147 , 1. ],
[ 5.2524 , -1.22 , 1. ],
[ 6.5894 , 5.9966 , 1. ],
[ 9.2482 , 12.134 , 1. ],
[ 5.8918 , 1.8495 , 1. ],
[ 8.2111 , 6.5426 , 1. ],
[ 7.9334 , 4.5623 , 1. ],
[ 8.0959 , 4.1164 , 1. ],
[ 5.6063 , 3.3928 , 1. ],
[12.836 , 10.117 , 1. ],
[ 6.3534 , 5.4974 , 1. ],
[ 5.4069 , 0.55657, 1. ],
[ 6.8825 , 3.9115 , 1. ],
[11.708 , 5.3854 , 1. ],
[ 5.7737 , 2.4406 , 1. ],
[ 7.8247 , 6.7318 , 1. ],
[ 7.0931 , 1.0463 , 1. ],
[ 5.0702 , 5.1337 , 1. ],
[ 5.8014 , 1.844 , 1. ],
```

```
[11.7 , 8.0043 , 1. ],  
[ 5.5416 , 1.0179 , 1. ],  
[ 7.5402 , 6.7504 , 1. ],  
[ 5.3077 , 1.8396 , 1. ],  
[ 7.4239 , 4.2885 , 1. ],  
[ 7.6031 , 4.9981 , 1. ],  
[ 6.3328 , 1.4233 , 1. ],  
[ 6.3589 , -1.4211 , 1. ],  
[ 6.2742 , 2.4756 , 1. ],  
[ 5.6397 , 4.6042 , 1. ],  
[ 9.3102 , 3.9624 , 1. ],  
[ 9.4536 , 5.4141 , 1. ],  
[ 8.8254 , 5.1694 , 1. ],  
[ 5.1793 , -0.74279, 1. ],  
[21.279 , 17.929 , 1. ],  
[14.908 , 12.054 , 1. ],  
[18.959 , 17.054 , 1. ],  
[ 7.2182 , 4.8852 , 1. ],  
[ 8.2951 , 5.7442 , 1. ],  
[10.236 , 7.7754 , 1. ],  
[ 5.4994 , 1.0173 , 1. ],  
[20.341 , 20.992 , 1. ],  
[10.136 , 6.6799 , 1. ],  
[ 7.3345 , 4.0259 , 1. ],  
[ 6.0062 , 1.2784 , 1. ],  
[ 7.2259 , 3.3411 , 1. ],  
[ 5.0269 , -2.6807 , 1. ],  
[ 6.5479 , 0.29678, 1. ],  
[ 7.5386 , 3.8845 , 1. ],  
[ 5.0365 , 5.7014 , 1. ],  
[10.274 , 6.7526 , 1. ],  
[ 5.1077 , 2.0576 , 1. ],  
[ 5.7292 , 0.47953, 1. ],  
[ 5.1884 , 0.20421, 1. ],  
[ 6.3557 , 0.67861, 1. ],  
[ 9.7687 , 7.5435 , 1. ],  
[ 6.5159 , 5.3436 , 1. ],  
[ 8.5172 , 4.2415 , 1. ],  
[ 9.1802 , 6.7981 , 1. ],  
[ 6.002 , 0.92695, 1. ],  
[ 5.5204 , 0.152 , 1. ],  
[ 5.0594 , 2.8214 , 1. ],  
[ 5.7077 , 1.8451 , 1. ],  
[ 7.6366 , 4.2959 , 1. ],  
[ 5.8707 , 7.2029 , 1. ],
```

```
[ 5.3054 ,  1.9869 ,  1.      ],
[ 8.2934 ,  0.14454,  1.      ],
[13.394  ,  9.0551 ,  1.      ],
[ 5.4369 ,  0.61705,  1.      ]])
```

```
In [155... data_val = data.values
m = len(data_val[:-1])
X = data[['x0', 'Population']].iloc[:-1].values
y = data['Profit'][:-1].values.reshape(m,1)
theta = np.zeros((2,1))

m, X.shape, y.shape, theta.shape
```

```
Out[155... (96, (96, 2), (96, 1), (2, 1))
```

```
In [156... theta
```

```
Out[156... array([[0.],
              [0.]])
```

```
In [157... computeCost(X,y,theta)
```

```
Out[157... 81.94398512635416
```

Gradient Descent

```
In [158... def gradientDescent(X,y,theta,alpha,num_iters):
    """
    Take numpy array for X,y,theta and update theta for every iteration of gradient steps
    Return theta and the list of cost of theta during each iteration
    """

    m = len(y)
    J_history = []
    for i in range(num_iters):
        predictions = X.dot(theta)
        error = np.dot(X.transpose(), (predictions-y))
```

```

        descent = alpha * 1/m * error
        theta -= descent
        J_history.append(computeCost(X,y,theta))

    return theta, J_history

```

```

In [159... theta, J_history = gradientDescent(X,y,theta,0.001, 2000)

```

```

In [160... print(f"h(x) = {str(round(theta[0,0],2))} + {str(round(theta[1,0],2))}x1")


```

h(x) = 2.88 + 0.76x1

```


In [161... from mpl_toolkits.mplot3d import Axes3D
#Generating values for theta0, theta1 and the resulting cost value
theta0_vals=np.linspace(-10,10,100)
theta1_vals=np.linspace(-1,4,100)
J_vals=np.zeros((len(theta0_vals),len(theta1_vals)))
for i in range(len(theta0_vals)):
    for j in range(len(theta1_vals)):
        t=np.array([theta0_vals[i],theta1_vals[j]])
        J_vals[i,j]=computeCost(X,y,t)
#Generating the surface plot
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
surf=ax.plot_surface(theta0_vals,theta1_vals,J_vals,cmap="coolwarm")
fig.colorbar(surf, shrink=0.5, aspect=5)
ax.set_xlabel("$\Theta_0$")
ax.set_ylabel("$\Theta_1$")
ax.set_zlabel("$J(\Theta)$")
#rotate for better angle
ax.view_init(30,120)

```




```
In [162... plt.plot(J_history)
plt.xlabel("Iteration")
plt.ylabel("$J(\Theta)$")
plt.title("Cost function using Gradient Descent")
```

Out[162... Text(0.5, 1.0, 'Cost function using Gradient Descent')




```
In [163... plt.scatter(data['Population'], data['Profit'])
x_value = [x for x in range(25)]
y_value = [x*theta[1] + theta[0] for x in x_value]
plt.plot(x_value, y_value, color = 'r')
plt.xticks(np.arange(5,30,step=5))
plt.yticks(np.arange(-5,30,step=5))
plt.xlabel('Population (in 10,000s)')
plt.ylabel('Profit (in 10,000$)')
plt.title('Profit vs Population')
```

Out[163... Text(0.5, 1.0, 'Profit vs Population')



```
In [164... def predict(x,theta):
    """
    Takes in numpy array x and theta and returns predicted value of y
    """
    predictions = np.dot(theta.transpose(),x)
    return predictions[0]
```

```
In [165... data.tail(1)
```

Out[165...

	Profit	Population	x0
96	5.4369	0.61705	1

```
In [166... predict1 = predict(data[['x0','Population']].iloc[-1].values, theta)*10000
print(f'For a population of 6170 the predicted profit is ${predict1}')
```

For a population of 6170 the predicted profit is \$33425.787565695806

Multivariate Regression

```
In [167... data2 = pd.read_csv('ex1data2.txt',header=None)
data2.head()
```

Out[167...

	0	1	2
0	2104	3	399900
1	1600	3	329900
2	2400	3	369000
3	1416	2	232000
4	3000	4	539900

```
In [168... data2.describe()
```

Out[168...


	0	1	2
count	47.000000	47.000000	47.000000
mean	2000.680851	3.170213	340412.659574
std	794.702354	0.760982	125039.899586
min	852.000000	1.000000	169900.000000
25%	1432.000000	3.000000	249900.000000

	0	1	2
50%	1888.000000	3.000000	299900.000000
75%	2269.000000	4.000000	384450.000000
max	4478.000000	5.000000	699900.000000

```
In [169... data2.columns = ['Size of House(Sq. feet)', 'No. of Rooms', "Price"]
```

```
In [170... plt.scatter(data2['Size of House(Sq. feet)'], data2['Price'])
plt.xticks(np.arange(800,4478,step=360))
plt.yticks(np.arange(100000,699900,step=70000))
plt.xlabel('Size (in Sq.Feet)')
plt.ylabel('Price (in $)')
plt.title('Size vs Price')
```


```
Out[170... Text(0.5, 1.0, 'Size vs Price')
```



```
In [171... plt.scatter(data2['No. of Rooms'], data2['Price'])
plt.xticks(np.arange(0,6,step=1))
```

```
plt.yticks(np.arange(100000,699900,step=70000))
plt.xlabel('No. of Rooms')
plt.ylabel('Price (in $)')
plt.title('No. of Rooms vs Price')
```

Out[171...] Text(0.5, 1.0, 'No. of Rooms vs Price')



```
In [172...] data2.columns = ['Size of the house (in square feet)', 'Number of bedrooms', 'Price of the house']
data2.head()
```

```
Out[172...]
   Size of the house (in square feet)  Number of bedrooms  Price of the house
0                                2104                      3          399900
1                                1600                      3          329900
2                                2400                      3          369000
3                                1416                      2          232000
4                                3000                      4          539900
```

```
In [173...] def normalize(dt):
```

```

df = dt.copy()
for col in df.columns:
    df[col] = (df[col]-df[col].mean())/df[col].std()
return df

#normalized

```

influence the result more due to its larger value. But this doesn't necessarily mean it is more important as a predictor. So we normalize the data to bring all the variables to the same range

```

In [174... normalized_data2 = normalize(data2)

```

```

In [175... normalized_data2['x0'] = 1

```

```

In [176... data_val = normalized_data2.values
m = len(data_val[:-1])
X = normalized_data2[['x0','Size of the house (in square feet)','Number of bedrooms']].iloc[:-1].values
y = normalized_data2['Price of the house'][:-1].values.reshape(m,1)
n = X.shape[1]
theta = np.zeros((n,1))

m, X.shape, y.shape, theta.shape

```

```

Out[176... (46, (46, 3), (46, 1), (3, 1))

```

```

In [177... computeCost(X,y,theta)

```

```

Out[177... 0.9858408807221255

```

```

In [178... thet, Jhist = gradientDescent(X,y,theta,0.01,1000)

```

```


In [179... print(f"h(x) = {str(round(thet[0,0],4))} + {str(round(thet[1,0],4))}x1 + {str(round(thet[2,0],4))}x2")

```

$$h(x) = -0.0014 + 0.8805x_1 + -0.0478x_2$$

```
In [180... plt.plot(J_history)
plt.xlabel("Iteration")
plt.ylabel("$J(\Theta)$")
plt.title("Cost function using Gradient Descent")
```

Out[180... Text(0.5, 1.0, 'Cost function using Gradient Descent')



```
In [181... X_req = normalized_data2[['x0', 'Size of the house (in square feet)', 'Number of bedrooms']].iloc[-1].values
X_req
```

Out[181... array([1. , -1.00374794, -0.22367519])

```
In [182... thet
```

Out[182... array([[-0.00135533],
[0.88054216],
[-0.04783476]])

```
In [183...
```

```
data2.tail(1)
```

```
Out[183...
```

	Size of the house (in square feet)	Number of bedrooms	Price of the house
46	1203	3	239500

```
In [184... normalized_data2.tail(1)
```

```
Out[184...
```

	Size of the house (in square feet)	Number of bedrooms	Price of the house	x0
46	-1.003748	-0.223675	-0.807044	1

```
In [185... def find_val(pred):  
    mul = data2["Price of the house"].std()  
    ad = data2["Price of the house"].mean()  
    return pred*mul+ad
```

```
In [186... X_req = normalized_data2[['x0', 'Size of the house (in square feet)', 'Number of bedrooms']].iloc[-1].values  
predict1 = predict(X_req, thet)  
print(find_val(predict1))
```

```
231065.48439645045
```