

Aryaman Gautam

J001

SKLEARN API

LINEAR REGRESSION:

Linear Regression fits a linear model with coefficients $w = (w_1, \dots, w_p)$ to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation.

It's parameters are:

- `fit_intercept` bool, default=True

Whether to calculate the intercept for this model. If False then no intercept is used for calculation.

- `copy_X` bool, default=True

If True, X will be copied; else, it may be overwritten

- `n_jobs` int, default=None

The number of jobs to use for the computation. This will only provide speedup for `n_targets > 1` and sufficient large problems.

- `positive` bool, default=False

When set to True, forces the coefficients to be positive. This option is only supported for dense arrays.

`LinearRegression(*, fit_intercept=True, normalize=False, copy_X=True, n_jobs=None, positive=False)`

Attribute:

`coef_array` of shape `(n_features,)` or `(n_targets, n_features)`

Estimated coefficients for the linear regression problem.

rank_int	Rank of matrix X. Only available when X is dense.
singular_array of shape (min(X, y),)	Singular values of X. Only available when X is dense.
intercept_float or array of shape (n_targets,)	Independent term in the linear model. Set to 0.0 if fit_intercept = False.

Methods

fit(X, y[, sample_weight])	Fit linear model.
get_params([deep])	Get parameters for this estimator.
predict(X)	Predict using the linear model.
score(X, y[, sample_weight])	Return the coefficient of determination of the prediction.
set_params(**params)	Set the parameters of this estimator

LOGISTIC REGRESSION:

It is used for predicting the categorical dependent variable using a given set of independent variables. Logistic regression predicts the output of a categorical dependent variable.

Fit(X,y)-fit the model according to the given training data

Predict(x)-predict class labels

Score(X,y)-returns mean accuracy on the given test data and label

Code:

```
sklearn.linear_model.LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)
```

RIDGE:

$$\|y - Xw\|^2 + \alpha \|w\|^2$$

This model solves a regression model where the loss function is the linear least squares function and regularization is given by the l2-norm. Also known as Ridge Regression or Tikhonov regularization. This estimator has built-in support for multi-variate regression (i.e., when y is a 2d-array of shape (n_samples, n_targets)).

Ridge regression penalizes the model based on the sum of squares of magnitude of the coefficients.

Alpha-Regularization strength; must be a positive float. Regularization improves the conditioning of the problem and reduces the variance of the estimates. Larger values specify stronger regularization.

It's parameters are:

- `alpha`{float, ndarray of shape (n_targets,)}, default=1.0
Regularization strength; must be a positive float. Larger values specify stronger regularization
- `fit_intercept`: bool, default=True
Whether to fit the intercept for this model
- `normalize`: bool, default=False
This parameter is ignored when `fit_intercept` is set to False.
- `copy_X`bool, default=True

If True, X will be copied; else, it may be overwritten

- `max_iterint`, default=None

Maximum number of iterations for conjugate gradient solver

- `tolfloat`, default=1e-3

Precision of the solution.

- `solver`{'auto', 'svd', 'cholesky', 'lsqr', 'sparse_cg', 'sag', 'saga'}, default='auto'

Solver to use in the computational routines. 'auto' chooses the solver automatically based on the type of data.

- `random_stateint`, RandomState instance, default=None

Used when `solver == 'sag' or 'saga'` to shuffle the data

`Ridge(alpha=1.0, *, fit_intercept=True, normalize=False, copy_X=True, max_iter=None, tol=0.001, solver='auto', random_state=None)`

Methods :

<code>fit(X, y[, sample_weight])</code>	Fit Ridge regression model.
<code>get_params([deep])</code>	Get parameters for this estimator.
<code>predict(X)</code>	Predict using the linear model.
<code>score(X, y[, sample_weight])</code>	Return the coefficient of determination of the prediction.
<code>set_params(**params)</code>	Set the parameters of this estimator.

LASSO :

LASSO regression penalizes the model based on the sum of magnitude of the coefficients.

Technically the Lasso model is optimizing the same objective function as the Elastic Net with `l1_ratio=1.0` (no L2 penalty).

```
Lasso(alpha=1.0, *, fit_intercept=True, normalize=False, precompute=False, copy_X=True,  
max_iter=1000, tol=0.0001, warm_start=False, positive=False, random_state=None,  
selection='cyclic')
```

Parameters:

- `alpha` float, default=1.0

Constant that multiplies the L1 term. Defaults to 1.0. `alpha = 0` is equivalent to an ordinary least square, solved by the `LinearRegression` object. For numerical reasons, using `alpha = 0` with the Lasso object is not advised

- `fit_intercept` bool, default=True

Whether to calculate the intercept for this model. If set to False, no intercept will be used in calculations (i.e. data is expected to be centered).

- `Normalize` bool, default=False

This parameter is ignored when `fit_intercept` is set to False. If True, the regressors X will be normalized before regression by subtracting the mean and dividing by the l2-norm.

- `Precompute` bool or array-like of shape (n_features, n_features), default=False

Whether to use a precomputed Gram matrix to speed up calculations. The Gram matrix can also be passed as argument. For sparse input this option is always False to preserve sparsity.

- `copy_X` bool, default=True

If True, X will be copied; else, it may be overwritten.

- `max_iter` int, default=1000

The maximum number of iterations.

- `Tol` float, default=1e-4

The tolerance for the optimization: if the updates are smaller than tol, the optimization code checks the dual gap for optimality and continues until it is smaller than tol.

- `warm_start` bool, default=False

When set to True, reuse the solution of the previous call to fit as initialization, otherwise, just erases the previous solution..

- `Positive` bool, default=False

When set to True, forces the coefficients to be positive.

- `random_state`int, RandomState instance, default=None

The seed of the pseudo random number generator that selects a random feature to update. Used when `selection == 'random'`.

Pass an int for reproducible output across multiple function calls.

- `selection`{'cyclic', 'random'}, default='cyclic'

If set to 'random', a random coefficient is updated every iteration rather than looping over features sequentially by default. This (setting to 'random') often leads to significantly faster convergence especially when `tol` is higher than $1e-4$